

Graph-Based Modeling, Scheduling, and Verification for Intersection Management of Intelligent Vehicles

CS637A: Embedded and Cyber Physical Systems


Fall 2020: Project Presentation

Ashwin Shenai (180156)

Kshitij Kabeer (180366)

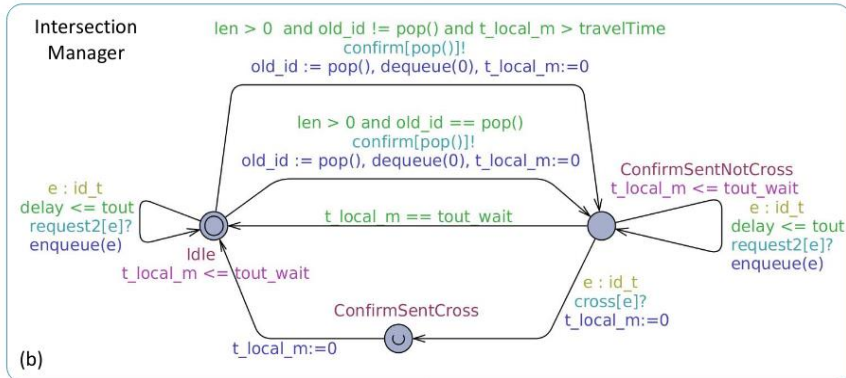
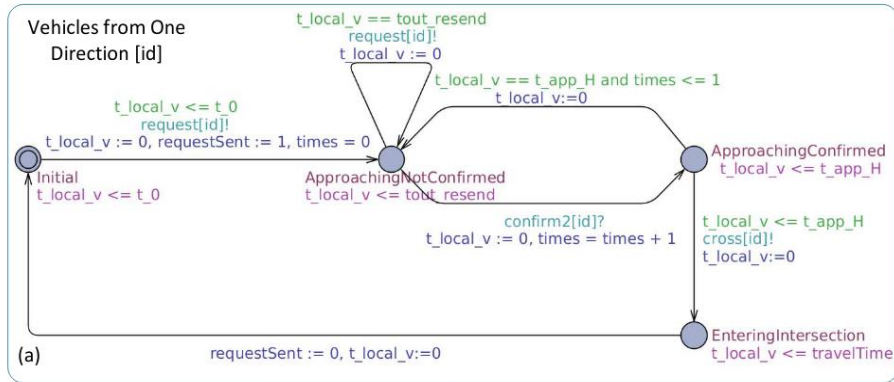
Intersection Management

- Management of vehicles and their passing order, at intersections
- Crucial for efficient traffic management and safety, especially with the advent of autonomous vehicles
- Optimizing passing time, preventing deadlock and ensuring no collisions – some of the prime objectives
- Position of each vehicle and commands communicated amongst themselves, or to a roadside unit – the intersection manager

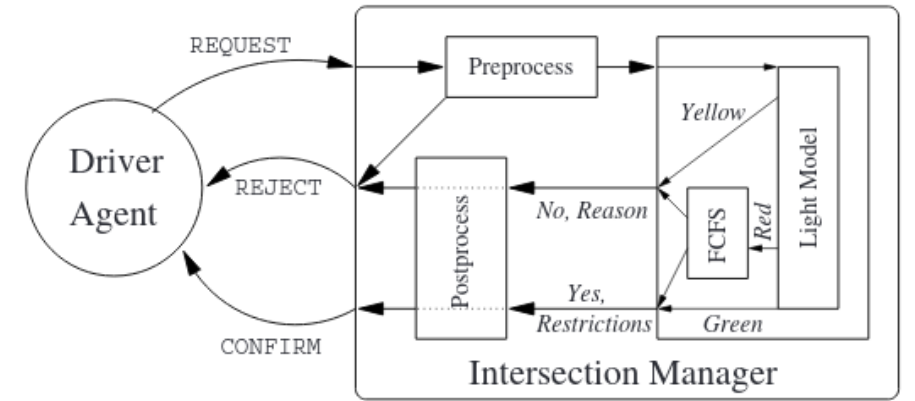





Related Work


- Protocols between vehicles and a centralized intersection manager



Delay-aware centralized intersection manager [26]



Multi-Agent Reservation-based Scheduler [12]



STIP: V2V Intersection Protocols [4]

Related Work


- Discrete-event control and conflict resolution in a centralized setting

Algorithm 1 Supervisor($\mathbf{x}(k\tau), \mathbf{u}_{driver}^k$)


```
1:  $\{\mathbf{T}_1, \mathbf{p}_1, answer_1\} = \text{Jobshop}(\hat{\mathbf{x}}(\mathbf{u}_{driver}^k), \Theta)$ 
2: if  $answer_1 = \text{yes}$  then
3:    $\mathbf{u}^{k+1, \infty} \leftarrow \sigma(\hat{\mathbf{x}}(\mathbf{u}_{driver}^k), \mathbf{T}_1, \mathbf{p}_1)$ 
4:    $\mathbf{u}_{safe}^{k+1} \leftarrow \mathbf{u}^{k+1, \infty}(t)$  for  $t \in [(k+1)\tau, (k+2)\tau)$ 
5:   return  $\mathbf{u}_{driver}^k$ 
6: else
7:    $\{\mathbf{T}_2, \mathbf{p}_2, answer_2\} = \text{Jobshop}(\hat{\mathbf{x}}(\mathbf{u}_{safe}^k), \Theta)$ 
8:    $\mathbf{u}^{k+1, \infty} \leftarrow \sigma(\hat{\mathbf{x}}(\mathbf{u}_{safe}^k), \mathbf{T}_2, \mathbf{p}_2)$ 
9:    $\mathbf{u}_{safe}^{k+1} \leftarrow \mathbf{u}^{k+1, \infty}(t)$  for  $t \in [(k+1)\tau, (k+2)\tau)$ 
10:  return  $\mathbf{u}_{safe}^k$ 
11: end if
```

Job scheduling-based semi-autonomous supervisory control

Refs: [2,7,9,22]



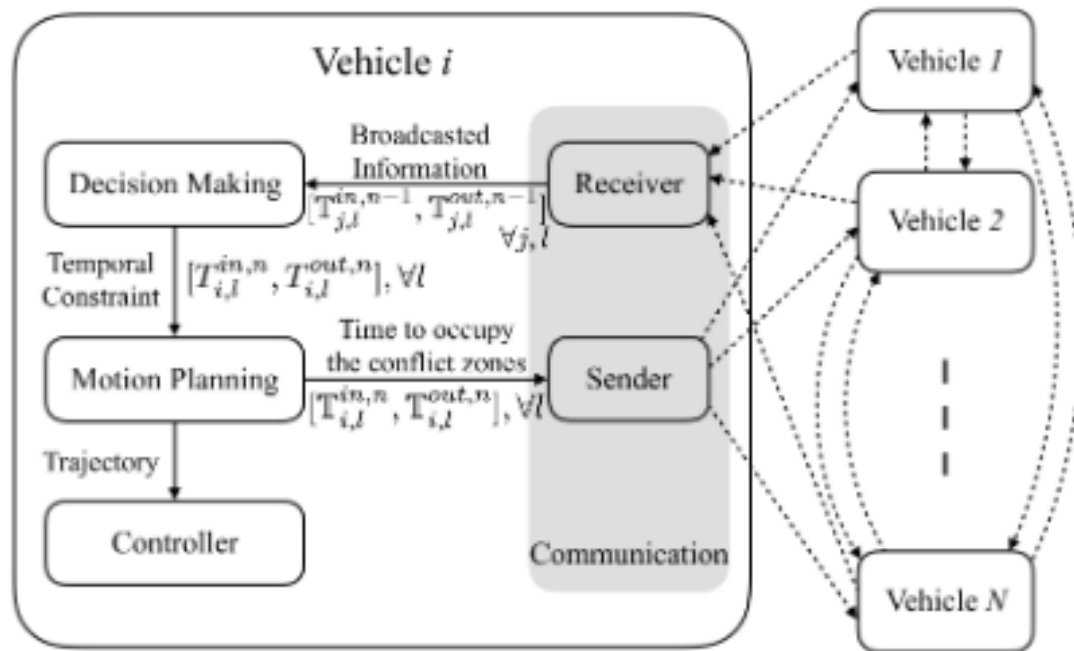
Discrete-event conflict-based modelling



Reactive supervisory control


Related Work

- Distributed inter-vehicle communication-based scheduling



Vehicle model for distributed scheduling [18]

- Petri net-based modelling for cooperative vehicles




Timed petri-net model for two-lane intersection [24]

Paper Contributions


- Graph based model – can deal with various granularities of intersections, highly expressive
- Centralized cycle removal for efficient, safe and deadlock free crossing of vehicle
- Efficiently scalable in response to increasing number of vehicles and conflict zone complexity
- Formal verification techniques to guarantee deadlock-freeness in all scenarios

Terminology




- Intersection
- Conflict Zone (j)
- Vehicle (Δ_i)
- Intersection Manager


- Earliest Arrival Time(a_i)




- Edge Waiting Time(w_k)



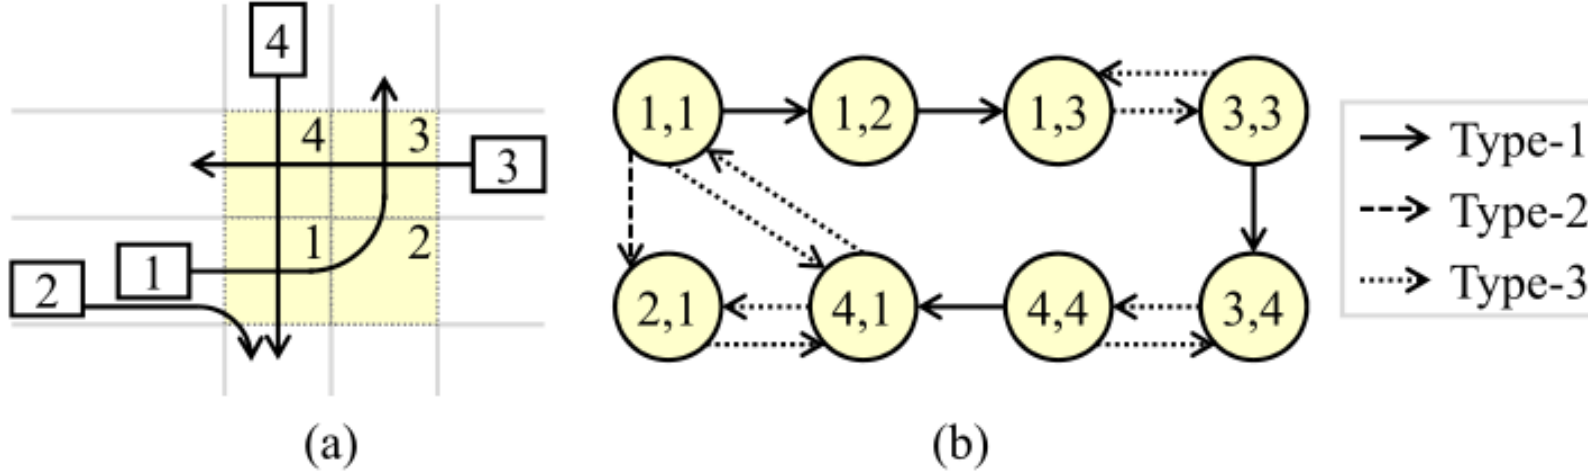
- Vertex Passing Time ($p_{i,j}$)



- Vertex Entering Time($s_{i,j}$)



Timing Conflict Graph (TCG)



- Type-1: Vehicle Δ_i goes from j to j'
- Type-2: Vehicles Δ_i and $\Delta_{i'}$ (in the same starting lane) go through j
- Type-3: Vehicles Δ_i and $\Delta_{i'}$ (in different starting lanes) go through j
 - Always in pairs

Problem Modelling

- Given a TCG G , earliest arrival times, edge waiting and passing times
1. Compute an acyclic subgraph G'
 - With all the vertices, Type-1 and Type-2 edges
 - Only one out of each pair of Type-3 edges
 2. Guarantee no deadlock in G'
 3. Assign an entering time to each vertex in G'
 4. Minimize the maximum leaving time $t_{\max} = \max_{G'}(s_{i,j} + p_{i,j})$
1. Collision Freeness
 2. Liveness/Feasibility
 3. Scheduling
 4. Optimality of Schedule

Assumptions

- Perfect, no-delay communication among vehicles and intersection managers.
 - Can model delay by increasing edge wait times, or adding noise in inputs.
- Problem solved in discrete chunks, no dynamic addition of vehicles
 - Vehicles coming in before the current graph is processed will be scheduled in the next chunk
- Dynamics of the vehicles aren't modelled – speed is constant or zero
 - No overtaking allowed

Verification


- Collision-freeness is guaranteed by the scheduler
- Need to ensure deadlock-freeness through verification
 - Graph-based verification
 - Petri net-based verification
- Either method can be used as a sub-routine to verify liveness of candidate schedules during scheduling

Graph-based Verification


- One would expect deadlock to occur when there is a cycle in the timing conflict graph. But:

Having no cycle in G' or G does not guarantee deadlock-freeness

- Deadlock can occur due to two parallel paths between same start and end vertices.
- Create an alternative graph to model deadlocks as cycles based on the timing conflict graph.



Ex. 1: V2 waits for V1 to pass Z1
V1 waits for V2 to pass Z3
Deadlock, no cycle



Ex. 2: V1 waits for V2 to pass Z1
V2 waits for V1 to pass Z3
No deadlock, V2 can move to Z2

Resource Conflict Graph (RCG)

- The basic idea is to combine edges of the conflict graph into vertices
- All Type-1 and Type-2 edges absorbed into vertices
- Each edge in the resource conflict graph is a Type-3 edge in the timing conflict graph
- At least one of the j -indices are equal across an edge

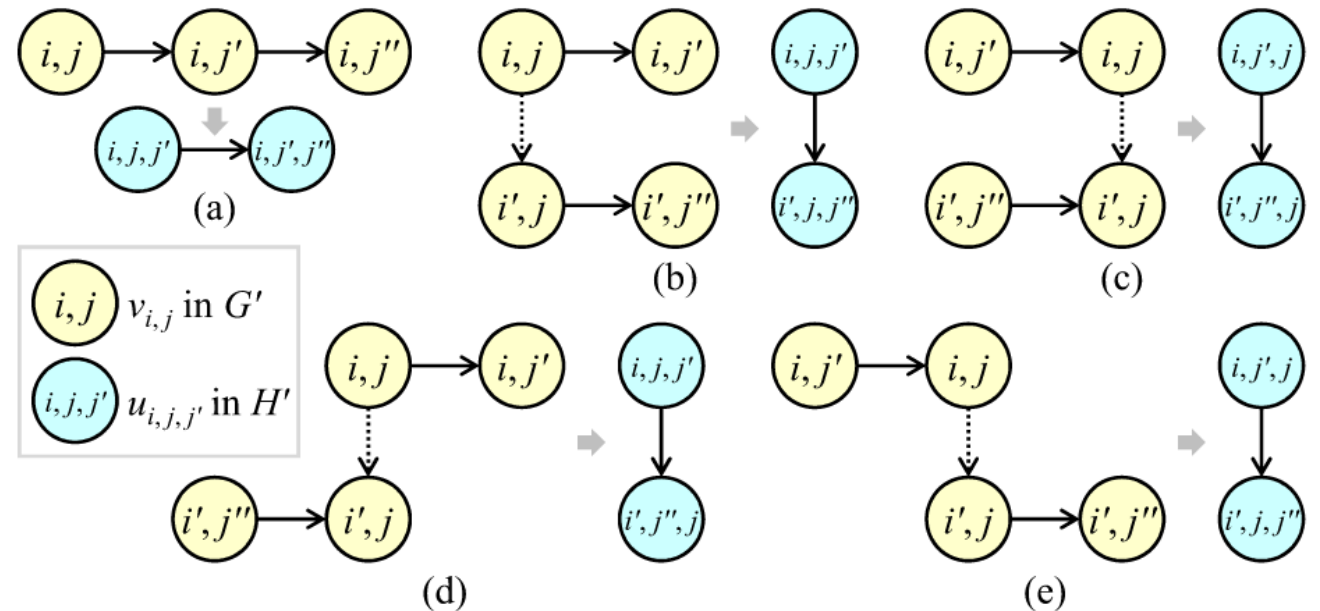




Fig. 6. The construction rules of resource conflict graphs.

Verifying Liveness

- An edge $(i_k, j_k, j'_k) \rightarrow (i_{k+1}, j_{k+1}, j'_{k+1})$ in RCG implies i_k must free up the common conflict zone before i_{k+1} arrives.
 - If there is a cycle in RCG, then there is a deadlock.
- If there is a deadlock, say i can't move from j to j' , then there must be an edge to (i, j, j') in RCG
 - Can show using one of the construction rules
 - Repeatedly apply above statement to all vehicles in deadlock
 - Ultimately forms a cycle in RCG, since vehicles and zones are finite.
- So, to verify that acyclic subgraph has no deadlock – construct its resource conflict graph and check for cycles in it.



Ex. 1: Deadlock, cycle exists




Ex. 2: No deadlock, no cycle


Petri Net Construction

Acyclic TCG G'



With Deadlock



Without Deadlock



Equivalent Petri-Net Π



Petri Net Verification

The Petri net Π has a deadlock if and only if G' has a deadlock

- If Π has a deadlock, at least one place $q_{i,i',j}$ never receives a token, which implies that Δ_i cannot leave j before $\Delta_{i'}$ enters j (so deadlock in G')
- If deadlock occurs in G' (suppose that some Δ_i can't go from j to j') it implies that $q_{i',i,j}$ will never receive a token (so deadlock in Π)
- So, to verify that acyclic subgraph has no deadlock – construct equivalent Petri Net and check it for deadlocks

Scheduling

- Naïve approach: first-come first-serve schedule
 - Ignores key interactions between vehicles and conflict zones
 - Introduces extra delays in many cases
- Generate a passing order for vehicles by constructing acyclic subgraph G' from conflict graph G with minimum total passing time.
 - Subgraph generated through cycle removal
- Naïve cycle removal: DFS traversal of the graph
 - May not always remove "good" edges to optimize objective
 - Can't remove some types of edges due to safety constraints
 - Minimum feedback arc set problem – intractable, no good approximations

Cycle Removal-Based Scheduling

- We need to remove cycles while minimizing total passing time
 - Min. Spanning Tree – acyclic subgraph with minimum sum of edge weights
 - Iteratively remove max-cost edge whose removal doesn't disconnect graph
- Proposed algorithm is based on the above idea
 - Iteratively remove max-cost Type-3 edges without violating constraints
 - Ensuring liveness complicates the problem – deadlocks exist even in acyclic graphs, as shown earlier
 - Need to efficiently handle cases where max-cost edge cannot be removed
 - Backtracking and redoing is computationally expensive.

Edge and Vertex States

Edge State: For an edge e ,

ON - e is included in G'

OFF – e has been removed from G'

UNDECIDED – Will decide ON/OFF in current subproblem

DONTCARE – e not included in current subproblem

- All Type-1 and Type-2 edges always ON

Vertex State: For a vertex v ,

BLACK – Entering time scheduled

GRAY – Entering time depends on Type-3 edges only

WHITE – Entering time can depend on any type of edges

- If any outgoing edge is ON, v is BLACK
- If v is BLACK, all edges through it must be ON/OFF
- If v is GRAY, v' must be BLACK if (v', v) is not a Type-3 edge

Vertex Entering Time

- Δ_i can't enter j before all earlier vehicles $\Delta_{i'}$ have passed

$$\max\{s_{i',j} + p_{i',j} + w_k\}$$

- Additionally, need to wait for $\Delta_{i'}$ to move to next zone j'

$$\max\{s_{i',j'} - w_{k'} + w_k\}$$

- Entering time is max of above two quantities – need to fulfill both
- For the first conflict zone on Δ_i 's path, also depends on arrival time a_i
- For v , depends on the earlier vertices u where (u,v) is an edge of G'
 - Since G' is acyclic, compute in topological order

Vertex Slack

- Maximum delay that can be added at vertex without changing the maximum leaving time t_{\max} (i.e. the optimization objective)
- For the last vertex on the path of a vehicle $v_{i,j'}$
$$t_{\max} - (s_{i,j'} + p_{i,j'})$$
- For other vertices u , it is minimum of slack of all reachable vertices v where (u,v) edge in G'
- Compute reverse topologically for acyclic graph

Insert example here

Defining the Cost of an Edge

Edge Cost: Delay incurred in t_{\max} due to adding this edge in G'

Only need to look at cost of Type-3 edges, $e_k = (v_{i,j}, v_{i',j})$

$$\text{cost}[e_k] = (s_{i,j} + p_{i,j} + w_k) - s_{i',j} - \text{slack}[v_{i',j}]$$

$(s_{i,j} + p_{i,j} + w_k)$ and $s_{i',j}$ are start times for $v_{i',j}$ with & without e_k

Compare with slack at $v_{i',j}$ to determine effect of e_k on t_{\max}

If the cost is positive, t_{\max} will increase.

But if cost is negative, t_{\max} won't change.

Removal of Type-3 Edges

- Initialization
 - Include Type-1 and Type-2 edges in G' , set their states to ON
 - Compute vertex entering times on G' , leaving time of last vehicle as t_{\max}
 - Set Type-3 states to UNDECIDED, compute vertex slacks.
- Identify candidate edges for removal
 - Leader vertex – $v_{i,j}$ where i is first vehicle on source lane, j is first conflict zone
 - Candidate edges – UNDECIDED Type-3 edges with one vertex as a leader vertex
 - Compute cost of these edges, and try to remove in decreasing order of cost

Ensuring Deadlock-Freeness

- Type-3 edges always in pairs – exactly one of two must be included
- Remove one and verify deadlock-freeness - if it fails, swap the edges
 - Use edge state variables to temporarily remove an edge
- If G' is deadlock-free, recompute vertex entering times and slacks
 - Identify newly set GRAY vertices as leader vertices, and repeat
- If G' is not deadlock-free, need to re-evaluate entire assignment till e_k
 - Backtracking is expensive – divide into subproblems
 - Schedule the first half of the vehicles arranged in increasing arrival time
 - For Type-3 edges between the two halves, assume first half passes before second half
 - Use the schedule of the first half while solving the second subproblem

Proof of Correctness and Time Complexity

- Type-1 and Type-2 edges included in G' by default.
- Exactly one Type-3 edge is selected out of every pair
- Deadlock-freeness is verified on removing each Type-3 edge
- For the solution obtained by recursively dividing into subproblems
 - No deadlock while merging both halves – we assume first passes before second
 - Each subproblem is essentially applying the same algorithm on a smaller set
 - Base case – only one vehicle: no Type-3 edges, so G' is feasible here
- Hence algorithm provably generates acyclic and deadlock-free G' .
- Time complexity of scheduling algorithm: $O(E^2 \log V)$

Results

- Paper results
- Our implementation results

Conclusion and References.